

An Overview of Machine Learning Models

Engr. Dr. Muhammad Farooq-i-Azam

Department of Computer Engineering
COMSATS University Islamabad, Lahore Campus

fazam@cuilahore.edu.pk

Introduction

- Machine Learning (ML) enables computers to learn patterns from data and make decisions or predictions.
- ML models are integral to applications like predictive text, ride-sharing ETAs, and content recommendations.
- This presentation explores key ML model categories, their mechanisms, and applications.

Supervised Learning

- Involves training models on labeled datasets, where each input is paired with a known output.
- The model learns to predict the output for new, unseen inputs.
- Common algorithms:
 - Linear Regression
 - Logistic Regression
 - Decision Trees
 - Support Vector Machines (SVMs)
 - Neural Networks
- Applications:
 - Email spam detection
 - Image recognition
 - Predictive analytics

Unsupervised Learning

- Deals with unlabeled data; the model seeks to identify inherent patterns or groupings.
- Types:
 - **Clustering:** This finds the natural groupings for all data.
 - **Association:** The dependencies or interesting relationships between various data are determined.
 - **Dimensionality Reduction:** Dimensions of data are reduced by finding the intrinsic components that represent certain data.
- Common algorithms:
 - K-Means Clustering
 - Hierarchical Clustering
 - Principal Component Analysis (PCA)
 - Autoencoders
- Applications:
 - Customer segmentation
 - Anomaly detection
 - Data compression

Semi-Supervised Learning

- Combines supervised and unsupervised learning by using a small amount of labeled data alongside a larger set of unlabeled data.
- The model leverages the labeled data to guide the learning process and improve accuracy.
- Applications:
 - Web content classification
 - Speech recognition
 - Protein sequence classification

Self-Supervised Learning

- A subset of unsupervised learning where the data itself provides the supervision.
- The model generates pseudo-labels from the input data and learns to predict these labels.
- Commonly used in:
 - Natural Language Processing (e.g., word embeddings)
 - Computer Vision (e.g., image colorization)
- Bridges the gap between unsupervised and supervised learning by creating supervisory signals from the data.

Reinforcement Learning

- Models learn by interacting with an environment, receiving rewards or penalties for actions taken.
- The goal is to develop a policy that maximizes cumulative rewards.
- Key components:
 - Agent: The learner or decision-maker.
 - Environment: Everything the agent interacts with.
 - Actions: Choices the agent can make.
 - Rewards: Feedback from the environment based on actions.
- Applications:
 - Game AI (e.g., AlphaGo)
 - Robotics control
 - Personalized recommendations

- **Type:** Model-free, off-policy algorithm.
- **Goal:** Learn the optimal action-selection policy.
- **Method:**
 - Learns Q-values: estimates of the total future rewards for taking an action in a given state.
 - Updates Q-values using the Bellman equation.
 - Chooses actions that maximize the Q-value, regardless of the current policy.
- **Advantages:**
 - Converges to optimal policy under certain conditions.
 - Simple to implement and widely used.
- **Use Case:** Grid world pathfinding, decision-making tasks.

SARSA (State-Action-Reward-State-Action)

- **Type:** Model-free, on-policy algorithm.
- **Goal:** Learn Q-values by following the current policy.
- **Method:**
 - Updates Q-values using the actual action taken under the current policy.
 - Update rule: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$
 - Accounts for the policy's behavior during learning.
- **Differences from Q-Learning:**
 - Q-Learning is off-policy (targets optimal action), SARSA is on-policy (uses actual action).
 - SARSA can be more cautious and stable in noisy environments.
- **Use Case:** Situations where risk-aware learning is important.

Temporal Difference (TD) Learning

- **Type:** Model-free prediction method.
- **Goal:** Estimate value functions by learning from incomplete episodes.
- **Method:**
 - Uses the difference (TD error) between predicted and actual rewards over time steps.
 - TD Update: $V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$
 - Balances the benefits of Monte Carlo methods and dynamic programming.
- **Advantages:**
 - Can learn online and incrementally.
 - Doesn't require waiting for final outcomes like Monte Carlo.
- **Use Case:** Policy evaluation in RL tasks.

Deep Q-Network (DQN)

- **Type:** Deep reinforcement learning algorithm.
- **Goal:** Extend Q-learning to handle high-dimensional state spaces.
- **Method:**
 - Uses a deep neural network to approximate the Q-function.
 - Inputs a state; outputs Q-values for each possible action.
 - Trained using experience replay and fixed target networks to improve stability.
- **Advantages:**
 - Works well on image and complex input data.
 - Capable of human-level performance on Atari games.
- **Use Case:** Game AI, robotics, autonomous systems.

Transfer Learning: Overview

- **Definition:** Transfer learning is a technique where knowledge from one task is reused to improve learning on a different, but related, task.
- **Why it Matters:**
 - Training deep learning models from scratch requires large datasets and computational resources.
 - Transfer learning allows leveraging pre-trained models to save time and improve performance.
- **Key Idea:** Use a model trained on a source task to boost performance on a target task.

- **Typical Workflow:**

- ① **Pretraining:** Train a model on a large, general-purpose dataset (e.g., ImageNet, Wikipedia).
- ② **Feature Reuse:** Keep the earlier layers (features) of the pretrained model.
- ③ **Fine-Tuning:** Replace and retrain the final layers on the target task dataset.

- **Layer Strategy:**

- Freeze low-level layers (generic features).
- Retrain high-level layers (task-specific features).

- **Applications:**

- **Computer Vision:** Using pretrained CNNs for tasks like facial recognition or medical imaging.
- **Natural Language Processing:** Adapting models like BERT, GPT for tasks like sentiment analysis, chatbots, summarization.

- **Benefits:**

- Reduces training time and computational cost.
- Requires less labeled data for the new task.
- Often improves performance on small or domain-specific datasets.

Deep Learning Architectures

- Deep Learning involves neural networks with multiple layers (deep neural networks) that can model complex patterns in data.
- Common architectures:
 - Convolutional Neural Networks (CNNs): Specialized for processing grid-like data such as images.
 - Recurrent Neural Networks (RNNs): Designed for sequential data, capturing temporal dependencies.
 - Transformers: Utilize self-attention mechanisms, excelling in tasks like language modeling.
- Applications:
 - Image and speech recognition
 - Natural language processing
 - Autonomous vehicles

- Combine predictions from multiple models to improve accuracy and robustness.
- Common techniques:
 - Bagging: Builds multiple independent models and averages their predictions (e.g., Random Forests).
 - Boosting: Builds models sequentially, each correcting errors of the previous one (e.g., Gradient Boosting Machines).
- Applications:
 - Competition-winning solutions in machine learning contests
 - Risk assessment in finance
 - Medical diagnosis

Conclusion

- Understanding various ML models is crucial for selecting the appropriate approach for a given problem.
- While foundational models like supervised and unsupervised learning are widely known, emerging techniques like self-supervised learning are gaining prominence.
- Advanced architectures and ensemble methods further enhance the capabilities of ML systems.
- Continuous learning and adaptation are key in the evolving field of machine learning.

- Christopher M. Bishop, Pattern Recognition and Machine Learning, Springer
- Will Jennings, ML Models: Understanding the Fundamentals, Gretel
- Stephen DeAngelis, Algorithms: Machine Learning's Secret Sauce, Enterra Solutions